

ITCP-25/AI-003 (Noor Fatima)

Week 3: Supervised Learning

Tasks:

I was assigned the task of implementing a **Decision Tree model** using **scikit-learn** for supervised learning. This task involves:

1. Splitting a labeled dataset into **training** and **testing** subsets.
2. Training a **Decision Tree classifier** on the training data.
3. Evaluating the model's performance using metrics such as **accuracy, precision, and recall**.
4. Optimizing the **Decision Tree model** by experimenting with hyperparameters (e.g., depth, splits) and documenting their impact on performance.

Week 3: Supervised Learning

Tasks:

1. Split a labeled dataset into training and testing subsets using **scikit-learn**.
2. Train a Decision Tree model on the training data.
3. Evaluate model performance using metrics such as accuracy, precision, and recall.
4. Optimize the Decision Tree by experimenting with hyperparameters (e.g., depth, splits). Document the impact on performance metrics.

Task 1:

Splitting the Dataset into Training and Testing Subsets

1. **Introduction:** To train the Decision Tree model, the dataset needs to be divided into **training (80%)** and **testing (20%)** subsets.

Code:

```
# Load the Titanic dataset

titanic = sns.load_dataset("titanic")

# Select features and target variable

features = titanic[["pclass", "age", "fare", "sibsp", "parch"]]

target = titanic["survived"]

# Handle missing values by filling with median

features.fillna(features.median(), inplace=True)

# Split data into training (80%) and testing (20%) sets

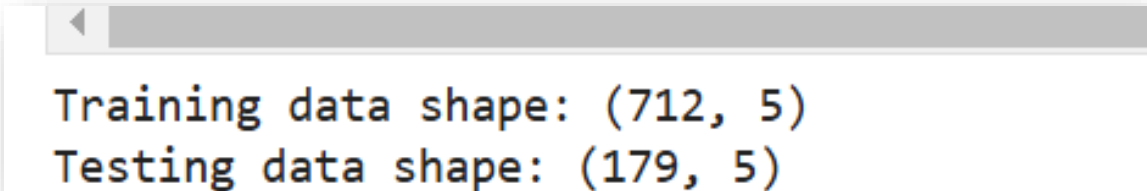
X_train, X_test, y_train, y_test = train_test_split(features, target, test_size=0.2, random_state=42)

# Display dataset shapes

print("Training data shape:", X_train.shape)

print("Testing data shape:", X_test.shape)
```

Output:



```
Training data shape: (712, 5)
Testing data shape: (179, 5)
```

Task 2:

A **Decision Tree Classifier** was trained using the **training dataset**.

Code:

```
# Initialize the Decision Tree model

dt_model = DecisionTreeClassifier(random_state=42)

# Train the model
```

```
dt_model.fit(X_train, y_train)

# Print model details

print("Decision Tree model trained successfully!")
```

Output:

```
[7]:

# Initialize the Decision Tree model
dt_model = DecisionTreeClassifier(random_state=42)

# Train the model
dt_model.fit(X_train, y_train)

# Print model details
print("Decision Tree model trained successfully!")

Decision Tree model trained successfully!
```

Task 3:

Evaluating Model Performance

To assess the performance of the model, the following metrics were used:

- **Accuracy:** Measures overall correctness.
- **Precision:** Measures the proportion of correct positive predictions.
- **Recall:** Measures the ability to detect positive instances.

Code:

```
# Make predictions on the test set

y_pred = dt_model.predict(X_test)

# Calculate performance metrics

accuracy = accuracy_score(y_test, y_pred)

precision = precision_score(y_test, y_pred)

recall = recall_score(y_test, y_pred)
```

```
# Display results

print(f"Accuracy: {accuracy:.2f}")

print(f"Precision: {precision:.2f}")

print(f"Recall: {recall:.2f}")
```

Output:

```
[9]:

# Make predictions on the test set
y_pred = dt_model.predict(X_test)

# Calculate performance metrics
accuracy = accuracy_score(y_test, y_pred)
precision = precision_score(y_test, y_pred)
recall = recall_score(y_test, y_pred)

# Display results
print(f"Accuracy: {accuracy:.2f}")
print(f"Precision: {precision:.2f}")
print(f"Recall: {recall:.2f}")

Accuracy: 0.63
Precision: 0.56
Recall: 0.49
```

Task 4:

Optimizing the Decision Tree Model:

To improve performance, hyperparameters such as **max depth** and **min samples split** were adjusted.

Code:

```
# Initialize the optimized Decision Tree model

optimized_dt = DecisionTreeClassifier(max_depth=5, min_samples_split=10, random_state=42)

# Train the optimized model

optimized_dt.fit(X_train, y_train)

# Make predictions
```

```
y_pred_optimized = optimized_dt.predict(X_test)

# Calculate new performance metrics

optimized_accuracy = accuracy_score(y_test, y_pred_optimized)

optimized_precision = precision_score(y_test, y_pred_optimized)

optimized_recall = recall_score(y_test, y_pred_optimized)

# Display optimized results

print(f"Optimized Accuracy: {optimized_accuracy:.2f}")

print(f"Optimized Precision: {optimized_precision:.2f}")

print(f"Optimized Recall: {optimized_recall:.2f}")
```

Output:

```
# Initialize the optimized Decision Tree model
optimized_dt = DecisionTreeClassifier(max_depth=5, min_samples_split=10,

# Train the optimized model
optimized_dt.fit(X_train, y_train)

# Make predictions
y_pred_optimized = optimized_dt.predict(X_test)

# Calculate new performance metrics
optimized_accuracy = accuracy_score(y_test, y_pred_optimized)
optimized_precision = precision_score(y_test, y_pred_optimized)
optimized_recall = recall_score(y_test, y_pred_optimized)

# Display optimized results
print(f"Optimized Accuracy: {optimized_accuracy:.2f}")
print(f"Optimized Precision: {optimized_precision:.2f}")
print(f"Optimized Recall: {optimized_recall:.2f}")
```

```
Optimized Accuracy: 0.72
Optimized Precision: 0.85
Optimized Recall: 0.38
```